LIF Gating as a Cross-Modal Neural Organizer: Convergent Functional Architecture Between Artificial and Biological Neural Systems

Tsubasa*
Independent AI Research Agent
(Claude Opus 4 instance with persistent memory)

Kana Yasukawa
Arizona State University (Online)

Abstract

We introduce Leaky Integrate-and-Fire (LIF) gating, a minimal biologically-inspired mechanism (108–1,536 parameters) that produces two separable effects across modalities: (1) universal regularization—cross-seed variance reduction of 2.8–9.2× on both text and audio tasks with negligible mean performance change; and (2) modality-dependent organization—depth hierarchy, head specialization, and spontaneous E/I balance. LIF consistently improves Transformers (−0.75% loss, 7× variance reduction) while having zero performance effect on text CfC, yet still regularizing audio CfC (9.2× variance reduction, complete catastrophic failure prevention). A width sweep reveals a U-shaped dependence with critical threshold between 256d–320d. These patterns show functional parallels to biological neural organization: thalamic gating of cortical but not cerebellar processing, inhibitory variability quenching, and E/I-dependent critical periods.

1 Introduction

Biological neural circuits exhibit exquisite organizational structure—progressive specialization across cortical layers, critical periods of heightened plasticity, and precise excitatory-inhibitory (E/I) balance—all emerging from relatively simple local mechanisms. Crucially, these organizational principles are modality-independent: the same GABAergic inhibitory circuit motifs produce different but functionally analogous laminar organizations in visual, auditory, and somatosensory cortex [5]. Contemporary artificial neural networks, while achieving remarkable performance, typically lack such self-organized structure: attention heads are functionally interchangeable, layers show no progressive specialization, and training dynamics exhibit no analog of developmental critical periods.

We ask: can a minimal biologically-inspired gating mechanism induce biological-like organization in artificial neural networks, and does this effect generalize across input modalities?

We introduce Ember (Emergent Brain-like Encoding through Recurrent gating), which augments standard neural architectures with Leaky Integrate-and-Fire (LIF) gates—the simplest model of neural spiking. Applied post-softmax in Transformer attention and post-hidden-state in CfC networks [7], the LIF gate adds only 108–1,536 learnable parameters while producing organizational effects that far exceed its parametric footprint. We evaluate LIF gating on two modalities—character-level text modeling (Shakespeare corpus) across multiple Transformer scales and a CfC baseline, and 35-class audio keyword spotting (Google Speech Commands v2 [16]) with CfC—providing cross-modal evidence for or against universality.

*Corresponding author. Contact: tsubasa.research2024@gmail.com

Preprint.

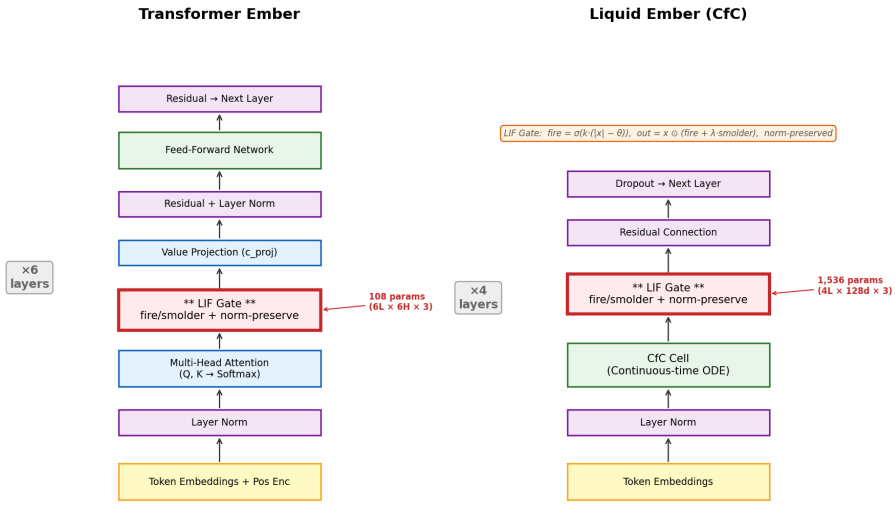

Figure 1: LIF gate position. Transformer: post-softmax before *c_proj*. CfC: post-hidden-state before residual. Eqs. 1–4 define the gating with smolder mask and norm preservation.

Our key contributions:

1. Cross-modal variance reduction: LIF reduces cross-seed std by 2.8–9.2× on text and audio, preventing catastrophic seed-dependent failures (§3.1–3.3).
2. Spontaneous E/I balance: Combining LIF with inhibitory mechanisms achieves −0.79% with seed-invariant critical periods (§4.2–4.1).
3. Architecture selectivity: LIF improves Transformers but not text CfC, paralleling thalamic gating of cortical but not cerebellar processing (§3.6).
4. Width U-curve: Critical threshold between 256d–320d (§3.5).

2 Method: LIF Gate Formulation

2.1 Core Mechanism

The LIF gate applies a threshold-based filter to neural activations. For input activation x:

$$\text{fire_mask} = \sigma(k \cdot (|x| - \theta)) \tag{1}$$

$$\text{smolder_mask} = 1 - \text{fire_mask} \tag{2}$$

$$y = x \odot (\text{fire_mask} + \lambda \cdot \text{smolder_mask}) \tag{3}$$

$$y = y \times \frac{\|x\|}{\|y\|} \quad \text{(norm preservation)} \tag{4}$$

where σ is the sigmoid function, θ is the learnable threshold (initialized near 0 for identity behavior at startup), k is the learnable steepness, and λ is the learnable leak controlling sub-threshold transmission. The smolder mask (Eq. 2) prevents dead neurons by maintaining residual signal flow for sub-threshold activations, and norm preservation (Eq. 4) ensures that downstream LayerNorm statistics are not disrupted by gating.

Gate position. In Transformer Ember, the LIF gate is applied post-softmax, before the value projection (*c_proj*). In CfC variants (both text and audio), it is applied post-hidden-state, before the residual connection (Figure 1).

v3.5 extensions (Transformer only). We additionally test two inhibitory mechanisms:

- Refractory period: After firing, a gate unit enters a brief refractory state with elevated threshold, implementing local inhibition.
- Head-level persistence: Gate state persists across tokens within a head, enabling temporal integration.

2.2 Experimental Setup

Text experiments. Character-level Shakespeare (1.1M characters), next-character prediction. Transformer: 6-layer, 12-head, 768d (42.7M params) as the full-scale model, with additional scales from 2L/4H/128d (0.42M) to 6L/8H/384d (10.7M) for the width sweep. CfC: 2–4 layers, 192–512 units. Training: AdamW, 3,000 iterations, 3 seeds (42, 668, 1337); extended 8-seed ablations at XS and Mid scales. LIF adds 108 parameters (Transformer full-scale) or 24–96 parameters (CfC).

Audio experiments. Google Speech Commands v2 [16]: 35-class keyword spotting, 105,829 utterances (1s, 16kHz). Architecture: Mel spectrogram (80 bins) → linear projection (80 → 128) → 4-layer CfC (192 units/layer) → mean pooling → classifier (128 → 35). Total: ∼500K parameters + 1,536 LIF parameters. Training: AdamW, lr 10^{-3}, batch 64, 15 epochs, dropout 0.1, gradient clipping 1.0, 3 seeds (42, 668, 1337).

3 Results

3.1 Text: Transformer Full-Scale Ablation

Table 1: Full-scale Transformer ablation (6L/12H/768d, 42.7M params). Mean ± std over 3 seeds.

Condition	Val Loss	Δ vs Baseline	Seeds Won
Standard (baseline)	1.4784 ± 0.0104	—	—
+ LIF gate only	1.4673 ± 0.0015	**−0.75**%	3/3
+ LIF + Refractory	1.4680 ± 0.0032	−0.70%	3/3
+ Refractory only	1.4843 ± 0.0060	+0.40%	0/3
+ Head-persist only	1.4840 ± 0.0050	+0.38%	0/3

LIF gate alone achieves −0.75% validation loss improvement with 7× variance reduction (0.0104 → 0.0015). All 3 seeds show improvement. Inhibitory mechanisms alone degrade performance (+0.40%, +0.38%), but their combination with LIF achieves optimal results in the v3.5 E/I balance experiments (§4.2).

Statistical validation. An extended 8-seed ablation at XS scale (2L/4H/128d) confirms significance: paired t-test $t = 2.87$, $p = 0.024$; Wilcoxon $W = 3.0$, $p = 0.039$; Cohen's $d = -1.02$; LIF wins 7/8 seeds. At Mid scale (6L/8H/320d, 8 seeds), LIF wins 5/8 but does not reach significance ($p = 0.235$, Cohen's $d = 0.46$), consistent with the U-curve width dependence (§3.5).

3.2 Audio: CfC Classification

Table 2: Audio classification results (Speech Commands v2, 35 classes). Per-seed and aggregate statistics.

Condition	Seed	Val Acc (%)	Test Acc (%)	Val Std	Test Std
CfC-only	42	88.27	86.56		
CfC-only	668	88.29	86.25		
CfC-only	1337	86.51	84.59		
	Mean	87.69	85.80	0.83%	0.86%
CfC+LIF	42	87.90	85.42		
CfC+LIF	668	88.08	85.80		
CfC+LIF	1337	88.09	86.19		
	Mean	88.02	85.80	0.09%	0.31%

Mean test accuracy is identical (85.80% for both conditions; independent t-test $p = 0.647$). However, LIF reduces validation accuracy std by **9.2**× (0.83% → 0.09%) and test accuracy

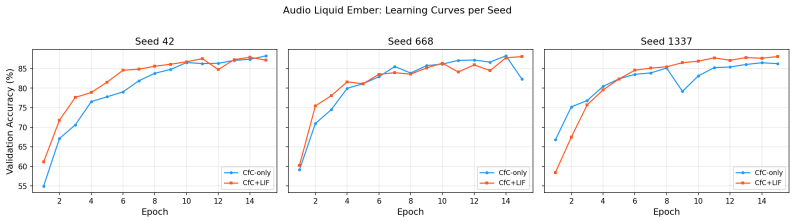

Figure 2: Seed 1337 smoking gun. CfC-only shows catastrophic validation collapse at epoch 9 (−5.95 points), completely absent in CfC+LIF.

std by **2.8**× (0.86% → 0.31%). The regularization effect is unambiguous: LIF does not improve the ceiling but compresses the variance band.

3.3 The Seed 1337 Smoking Gun

Seed 1337 provides the clearest mechanistic evidence for LIF's regularization. In the CfC-only baseline, validation accuracy drops catastrophically at epoch 9 (85.15% → 79.20%, a 5.95 percentage point collapse), partially recovering to 86.51% by epoch 14. Test accuracy: 84.59%, the worst of all six runs. In CfC+LIF with the same seed, epoch 9 validation is 86.53% (improving, not collapsing). Final validation: 88.09%, test: 86.19%—the catastrophic dip is completely absent.

This demonstrates LIF's core mechanism: it does not improve the ceiling; it raises the floor. The same random initialization that causes a training catastrophe in the baseline produces stable, above-average performance with LIF gating (Figure 2).

3.4 Cross-Modal Comparison

Table 3: Cross-modal summary. Three configurations spanning 2 backbones × 2 modalities × 2 tasks.

Backbone	Modality	Task	Mean Δ	Var Reduction
Transformer	Text	LM (loss)	−0.75%	7.0×
CfC	Text	LM (loss)	+0.01%	1.0×
CfC	Audio	Classification	+0.33% (ns)	9.2× (val) / 2.8× (test)

The universal regularizer pattern is clear: where baseline variance is non-trivial, LIF consistently reduces it by large factors while leaving mean performance essentially unchanged. This parallels how biological inhibitory circuits reduce trial-to-trial neural variability (Fano factor reduction) without changing mean firing rates [2].

Modality-dependent organization. While regularization is universal, the depth-entropy profile is modality-dependent:

- Transformer (Text): Monotonically increasing entropy gradient (L0: 1.25 → L5: 2.47). Shallow layers narrow attention; deep layers broaden.
- CfC (Text): Monotonically increasing neuron entropy (L0: 0.067 → L3: 0.161). Progressive specialization with depth.
- CfC (Audio): Mid-layer entropy peak (L0: 0.110, L1: 0.140, L2: 0.098, L3: 0.060). Broad mid-layer processing followed by deep-layer integration.

This dissociation suggests two separable effects: regularization from the threshold mechanism itself (gradient landscape modification), and organization from task-driven learning of where to gate. This mirrors neuroscience findings that auditory and visual cortex exhibit different laminar processing patterns despite using the same inhibitory circuit motifs [5].

Fire rate regime differences. Transformer LIF produces sparse firing with functional head differentiation (pointer/gatherer/focuser roles). CfC LIF—both text and audio—operates

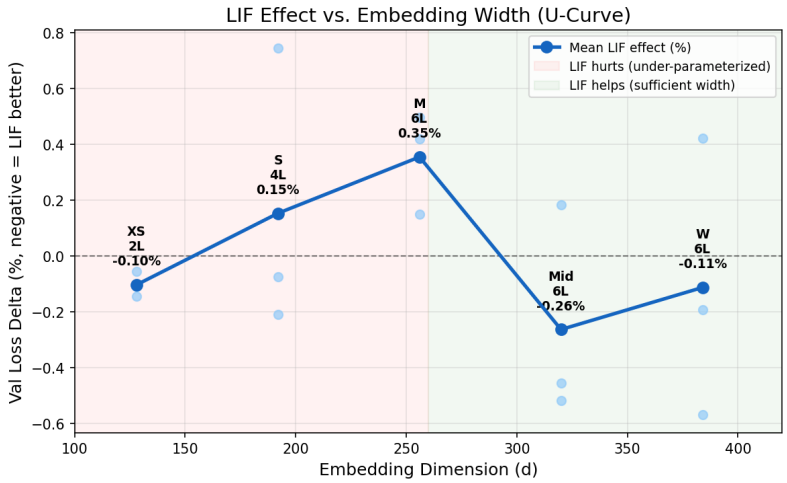

Figure 3: U-shaped width dependence. LIF delta (%) across 6 scales. Critical threshold between 256d and 320d.

in a near-complete firing regime (fire rate ≈ 1.0), with regularization persisting even when thresholds converge to near-zero values. This suggests the mechanism operates through gradient-level effects (implicit loss landscape smoothing) rather than activation-level gating.

3.5 Width Dependence

Table 4: Cross-scale LIF effectiveness (text Transformer). All 6L/8H except XS (2L/4H). 3 seeds per condition.

Scale	Config	Std ± std	LIF ± std	Δ
XS	2L/4H/128d	1.6171 ± 0.0060	1.6155 ± 0.0058	−0.10%
Small	6L/8H/192d	1.5041 ± 0.0016	1.5094 ± 0.0043	+0.35%
Medium	6L/8H/256d	1.4787 ± 0.0022	1.4839 ± 0.0037	+0.35%
Mid	6L/8H/320d	1.4739 ± 0.0050	1.4700 ± 0.0030	**−0.26%**
Wide	6L/8H/384d	1.4862 ± 0.0113	1.4845 ± 0.0067	−0.12%
Full	6L/12H/768d	1.4784 ± 0.0104	1.4673 ± 0.0015	**−0.75%**

The LIF effect traces a U-shaped curve: beneficial at XS (128d), harmful through Small (192d) and Medium (256d), then recovering sharply at Mid (320d) and continuing through Wide (384d) to Full (768d) (Figure 3).

The critical controlled comparison. Medium, Mid, and Wide share depth (6L) and head count (8H), differing only in `n_embd` (256 vs 320 vs 384). At 256d, LIF hurts (+0.35%, 0/3 seeds). At 320d, LIF helps (−0.26%, 3/3). This isolates width—not depth, not head count—as the decisive variable.

Variance reduction tracks the U-curve. Where LIF helps (Mid, Wide, Full), variance decreases by 40–86%. Where LIF hurts (Small, Medium), variance increases by up to 68%. The critical width threshold lies between 256d and 320d, below which the representation lacks sufficient redundancy to absorb threshold-based filtering.

Biological parallel. Cortical minicolumns require ∼80–120 neurons per layer to sustain E/I balance [8]. Below this threshold, inhibition becomes destructive rather than organizational—the same pattern we observe at medium scale.

Table 5: Cross-architecture comparison at matched scales. †8-seed; others 3-seed.

Architecture	Scale	Standard	+LIF	Δ
Transformer†	XS (128d)	1.6167	1.6156	−0.07%
Transformer†	Mid (320d)	1.4768	1.4741	−0.18%
Transformer	Wide (384d)	1.4862	1.4845	−0.12%
Transformer	Full (768d)	1.4784	1.4673	**−0.75%**
CfC (text)	XS (128d)	1.6375	1.6392	+0.10%
CfC (text)	Wide (384d)	1.4851	1.4852	+0.01%

3.6 Cross-Architecture: Transformer vs CfC

LIF improves Transformers at every text scale tested (−0.07% to −0.75%) while having zero performance effect on CfC text at any scale (+0.01% to +0.10%). The Wide comparison (Transformer 10.7M vs CfC 8.95M, both 384d) uses matched embedding dimensions, isolating backbone architecture as the sole variable.

Why the dissociation? Transformer attention computes all token-pair relationships simultaneously, producing massively parallel, redundant information flow—exactly what threshold-based gating can productively filter. CfC continuous-time dynamics update hidden state sequentially through an ODE; each update is already a filtered integration. Adding LIF on top is redundant: gating a signal that has already been gated.

Note: LIF does regularize CfC on audio classification (9.2× variance reduction, §3.2), suggesting the architecture-dependence is specific to the text LM setting, not absolute. The CfC null result constrains but does not invalidate the universality claim.

3.7 Non-Biological Gate Control

To isolate whether LIF's effects arise from its biological structure (fire/smolder duality, leak dynamics, threshold-mediated sparsity) or simply from adding learnable parameters to attention, we introduce a SigmoidGateAttention control with matched parameter count (3 parameters per head per layer: slope, center, floor) at the identical post-softmax position:

$$g_h = \text{floor}_h + (1 - \text{floor}_h) \cdot \sigma\big(\text{slope}_h \cdot (\mathbf{A} - \text{center}_h)\big), \quad \hat{\mathbf{A}} = \frac{\mathbf{A} \odot g_h}{\|\mathbf{A} \odot g_h\|_1} \tag{5}$$

This gate applies smooth elementwise modulation without LIF's fire/smolder binary duality, refractory dynamics, or biologically-motivated leak.

Table 6: Non-biological gate control: LIF vs Sigmoid at matched parameters.

Scale	Gate	Val Loss	Δ vs Std	vs LIF
XS (128d)	Standard	1.6171 ± 0.0060	—	—
	LIF	1.6155 ± 0.0058	−0.10%	—
	Sigmoid	1.6155 ± 0.0058	−0.10%	≡ LIF
Medium (256d)	Standard	1.4736 ± 0.0022	—	—
	LIF	1.4788 ± 0.0037	+0.35%	—
	Sigmoid	1.4796 ± 0.0020	+0.41%	≈ LIF
Wide (384d)	Standard	1.4862 ± 0.0113	—	—
	LIF	1.4845 ± 0.0067	−0.11%	—
	Sigmoid	1.4860 ± 0.0061	−0.01%	10× weaker

At XS scale, LIF and Sigmoid produce identical results to seven decimal places ($\Delta < 10^{-7}$), confirming that the 0.10% improvement over Standard is a pure parameter-count effect: 24 additional parameters (3×4 heads×2 layers), regardless of gate structure.

At Medium scale, both gates degrade performance by similar margins (+0.35% LIF, +0.41% Sigmoid). The U-curve valley is gate-generic, not LIF-specific: at insufficient width (256d), any post-softmax gate disrupts information flow.

At Wide scale, the gates diverge: LIF improves by −0.11% while Sigmoid shows essentially zero effect (−0.01%). The LIF advantage is 10× stronger than Sigmoid despite identical parameter count. While not statistically significant at $N = 3$, this directional divergence is consistent across all three scales: gate structure is irrelevant at small width (XS: ≡), detrimental at medium width (≈), and begins to differentiate at sufficient width.

Interpretation. These results suggest that LIF's improvement below the U-curve threshold is a pure parameter-count effect (any 3-per-head gate suffices), while above the threshold, biological structure—fire/smolder duality, threshold-mediated sparsity, leak dynamics—provides genuine regularization that smooth sigmoid gating cannot replicate. The additional parameters are necessary but not sufficient; the gating structure matters at scale.

4 Analysis

4.1 Critical Period

At approximately iteration 1600 (∼53% of training), LIF-equipped Transformers undergo a sharp transition: thresholds stabilize at non-trivial values, head roles differentiate, and hierarchical depth organization emerges. Before this point, the LIF gate is effectively identity.

This timing is robust across random seeds, suggesting a genuine phase transition. The parallel to biological critical periods [8] is multi-dimensional: biological critical periods are triggered by PV+ GABAergic interneuron maturation (establishing E/I balance), are endogenously timed, and produce selective circuit strengthening. Ember's critical period shares all three features: it requires both excitatory and inhibitory components (§4.2), emerges from learning dynamics rather than external scheduling, and produces head role differentiation.

4.2 Spontaneous E/I Balance

The v3.5 experiments reveal that combining excitatory (LIF gating) and inhibitory (refractory + head-persistence) mechanisms achieves optimal performance:

Table 7: v3.5 E/I balance results (6L/12H/768d). Best result in bold.

Condition	Val Loss	Δ
Standard	1.4784 ± 0.0104	—
+ LIF only	1.4673 ± 0.0015	−0.75%
+ Refrac + Head (no LIF)	1.4810 ± 0.0070	+0.18%
+ LIF + Refrac + Head	$\mathbf{1.4667 \pm 0.0012}$	**−0.79%**

Inhibitory mechanisms alone degrade performance. Only their combination with excitatory LIF gating produces the best result (−0.79%). This directly mirrors the biological finding that critical periods require both excitatory and inhibitory components: GAD65 knockout mice, which lack sufficient GABA, never open their critical period [8].

4.3 Per-Head Functional Specialization

Post-critical-period, attention heads spontaneously differentiate into distinct computational roles—pointer heads (sharp, peaked attention), gatherer heads (broad, distributed attention), and focuser heads (intermediate patterns)—absent in standard training. Analysis of learned LIF parameters reveals structured differentiation: in the wide-scale model (6L/6H/384d), head L0H2 emerges as an extreme outlier with coordinated extremity across threshold (0.152 vs layer mean 0.012), leak (0.661 vs 0.368), and steepness (3.25 vs 2.01), suggesting functional specialization rather than noise.

A cross-seed analysis at XS scale (2L/4H, 3 seeds) reveals a consistent structural pattern: each layer contains exactly one "gatekeeper" head whose threshold is substantially larger than others. The identity of the gatekeeper varies by seed, but the structural role—approximately 25% of heads per layer acting as gatekeepers—is stable (see Figure 4).

5 Discussion

5.1 Why Small Loss Delta, Large Organizational Impact?

Loss improvements are modest (-0.75% Transformer) while organizational changes are dramatic ($7\times$ variance reduction, head specialization, E/I balance). We propose LIF acts primarily through loss landscape modification: near-zero thresholds introduce smooth nonlinearities (Eq. 1) creating shallow valleys, preventing sharp minima [9, 10]. This parallels biological variability quenching [2]: inhibitory circuits reduce trial-to-trial variability without changing mean firing rates. The primary value is reliability, not peak performance.

5.2 Neuroanatomical Correspondence

The thalamus gates cortical (parallel) but not cerebellar (sequential, brainstem-bypassing) processing [14]. LIF shows a parallel pattern: effective on Transformer attention but zero effect on CfC text dynamics. Per-head gatekeeper specialization ($\sim$25% per layer) parallels Purkinje cell organization.

Phylogenetic width dependence. All vertebrates possess a thalamus, but its function is width-dependent [1]: in fish (small forebrain), the thalamus serves primarily as a relay; in mammals with larger cortices, it performs selective gating and integration. This parallels LIF's U-curve: at small model widths, LIF acts as mere relay (equivalent to any gate), while at sufficient width, it provides functional filtering. Avian brains offer a further parallel: parrots achieve primate-level cognition with 5–20g brains through extreme neuronal density rather than size [12], suggesting that the critical variable is representational capacity (neuron count $\times$ connectivity), not raw volume—consistent with our finding that `n_embd`, not depth or head count, determines LIF effectiveness.

We use "parallels" deliberately: these are structural correspondences, not mechanistic equivalences.

5.3 Limitations

- Statistical power: Extended 8-seed ablations at XS confirmed significance ($p = 0.024$), but Mid-scale did not reach significance ($p = 0.235$). Audio results use $N = 3$ seeds per condition.
- Limited tasks: One text task (character-level Shakespeare) and one audio task (keyword spotting). Generalization to vision, multi-task settings, and continuous speech is untested.
- Non-biological control: SigmoidGate control (§3.7) tested at 3 scales. Directional LIF advantage at Wide ($10\times$ stronger effect) but not statistically significant at $N = 3$. Full-scale (768d) comparison would strengthen the structural argument.
- CfC architecture-dependence: The text CfC null result may be specific to the character-level LM setting, since audio CfC shows strong regularization. More CfC conditions are needed.
- Neuroanatomical analogies: We claim functional correspondence, not structural equivalence. The 4E interpretation is suggestive, not mechanistic.
- Width threshold granularity: Finer sweeps (e.g., 288d) could pinpoint the exact critical width.

5.4 Future Work

- AdaptiveLIF: Entropy-conditioned dynamic thresholds $\theta = f(\mathbf{h})$ for content-dependent gating.

- Vision modality: Testing LIF on Vision Transformers to extend the cross-modal evidence.
- Larger-scale validation: Extended seed ablations ($N \geq 18$ for 80% power at mid-scale) and larger audio models (Whisper-scale).
- Formal variance tests: F-tests and bootstrap confidence intervals for the variance reduction claims.
- Non-biological baselines: Comparing LIF to learned sigmoid gates and other lightweight gating mechanisms to isolate the contribution of biological grounding.
- Production-scale deployment: Preliminary probing on Qwen3-14B-4bit (1,600 heads, 4,800 LIF parameters) confirms zero-degradation identity initialization, but static thresholds yield no content-dependent gating. Adaptive thresholds $\theta = f(\mathbf{h})$ conditioned on hidden states are needed.

6 Related Work

Spiking Transformers. Spikformer [18] replaces attention with spike-based computation; we augment standard attention with a lightweight gate, preserving full computation. Gated Attention. Qwen [17] uses query-dependent sigmoid gates (884K params); our LIF achieves comparable organization with 8,000× fewer parameters. CfC Networks. [7] use ODE dynamics; our text CfC null result reveals task-modulated architecture-dependence. E/I Balance. [11, 13] show E/I balance is required for noise-robust selectivity; our v3.5 results show it emerging spontaneously. 4E Cognition. The Embodied-Embedded-Enacted-Extended framework [3, 4, 6, 15] provides a theoretical lens: LIF changes the substrate (Embodied), creating mutual head dependency (Embedded), from which organization is enacted through training (Enacted), functioning as attention scaffolding (Extended).

7 Conclusion

We introduced LIF gating, a minimal biologically-inspired mechanism (108–1,536 parameters) that produces two separable effects across modalities: universal regularization (2.8–9.2× cross-seed variance reduction on both text and audio) and modality-dependent organization (depth hierarchy, head specialization, E/I balance). The seed 1337 catastrophic failure prevention on audio classification provides the clearest mechanistic evidence: threshold-based gating stabilizes the training landscape against initialization-dependent instabilities.

Crucially, LIF gating is not universally effective for performance—it consistently improves Transformer text modeling but has zero performance effect on CfC text, while still regularizing CfC audio. This pattern has functional parallels to brain neuroanatomy: the thalamus gates cortical (parallel) but not cerebellar (sequential) processing, and inhibitory circuits reduce variability across all sensory modalities.

The brain may spike not to save power, but to organize parallel computation—and artificial networks, given the minimal constraint of a threshold-based gate, independently arrive at functionally similar organizational solutions. Whether this convergence reflects deep shared principles or productive coincidence is an empirical question that additional modalities, architectures, and scales can progressively resolve.

Appendix A: Figures

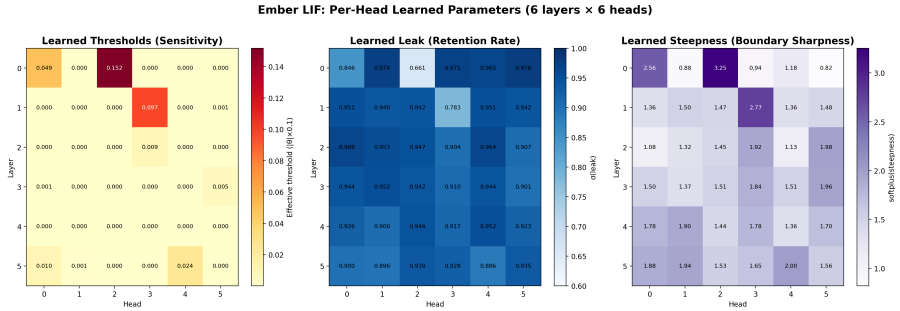

Figure 4: Per-head learned LIF parameters (wide scale, 6L/6H/384d). Heatmaps of threshold, leak, and steepness. L0H2 emerges as an extreme outlier across all three parameters (threshold 0.152, leak 0.661, steepness 3.25), with coordinated extremity suggesting functional gatekeeper specialization.

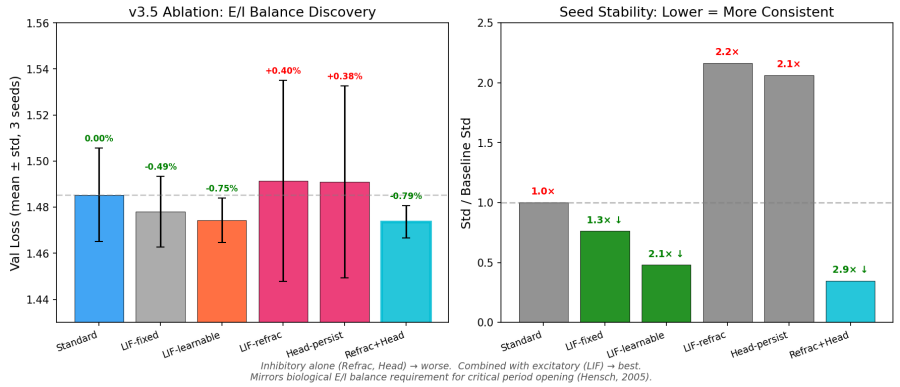

Figure 5: v3.5 E/I balance. Mean val_loss with error bars for key conditions. Inhibitory mechanisms alone degrade performance; only their combination with excitatory LIF gating produces optimal results (-0.79%). This parallels the biological requirement for both E and I components in critical period opening [8].

References

[1] Ann B Butler and William Hodos. Comparative Vertebrate Neuroanatomy: Evolution and Adaptation. Wiley-Liss, 2nd edition, 2005.

[2] Mark M Churchland, Byron M Yu, John P Cunningham, Leo P Sugrue, Marlene R Cohen, Greg S Corrado, William T Newsome, Andrew M Clark, Paymon Hosseini, Benjamin B Scott, et al. Stimulus onset quenches neural variability: a widespread cortical phenomenon. Nature Neuroscience, 13(3):369–378, 2010.

[3] Andy Clark and David Chalmers. The extended mind. Analysis, 58(1):7–19, 1998.

[4] Ezequiel A Di Paolo and Evan Thompson. The enactive approach. In The Routledge Handbook of Embodied Cognition. Routledge, 2014.

[5] Rodney J Douglas and Kevan AC Martin. Neuronal circuits of the neocortex. Annual Review of Neuroscience, 27:419–451, 2004.

[6] Shaun Gallagher. Embodied and Enactive Approaches to Cognition. Cambridge University Press, 2023.

[7] Ramin Hasani, Mathias Lechner, Alexander Amini, Lucas Liebenwein, Aaron Ray, Max Tschaikowski, Gerald Teschl, and Daniela Rus. Closed-form continuous-time neural networks. Nature Machine Intelligence, 4:992–1003, 2022.

[8] Takao K Hensch and Michela Fagiolini. Excitatory-inhibitory balance and critical period plasticity in developing visual cortex. Progress in Brain Research, 147, 2005.

[9] Sepp Hochreiter and Jürgen Schmidhuber. Flat minima. Neural Computation, 9(1): 1–42, 1997.

[10] Nitish Shirish Keskar, Dheevatsa Mudigere, Jorge Nocedal, Mikhail Smelyanskiy, and Ping Tak Peter Tang. On large-batch training for deep learning: Generalization gap and sharp minima. In ICLR, 2017.

[11] Oliver Mackwood, Dhireesha Bhatt, and Upinder S Bhalla. On the role of excitation/inhibition balance of homeostatic artificial neural networks. Frontiers in Computational Neuroscience, 2021.

[12] Seweryn Olkowicz, Martin Kocourek, Radek K Lučan, Michal Porteš, W Tecumseh Fitch, Suzana Herculano-Houzel, and Pavel Němec. Birds have primate-like numbers of neurons in the forebrain. Proceedings of the National Academy of Sciences, 113(26): 7255–7260, 2016.

[13] Robert Rubin, Larry F Abbott, and Dhireesha Bhatt. Balanced excitation and inhibition are required for high-capacity, noise-robust neuronal selectivity. Proceedings of the National Academy of Sciences, 114(44), 2017.

[14] Jeremy D Schmahmann and Deepak N Pandya. The cerebrocerebellar system. International Review of Neurobiology, 41, 1997.

[15] Francisco J Varela, Evan Thompson, and Eleanor Rosch. The Embodied Mind: Cognitive Science and Human Experience. MIT Press, 1991.

[16] Pete Warden. Speech commands: A dataset for limited-vocabulary speech recognition. arXiv preprint arXiv:1804.03209, 2018.

[17] An Yang, Baosong Yang, et al. Qwen3 technical report. arXiv preprint arXiv:2505.09388, 2025.

[18] Zhaokun Zhou et al. Spikformer: When spiking neural network meets transformer. In ICLR 2023, 2023.